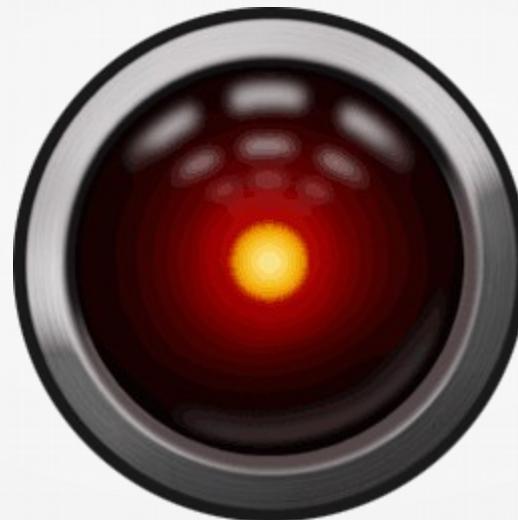


# Image/Video capture with V4L2



# Who am I



<https://github.com/Tropicao>

- Alexis Lothoré, embedded software engineer
  - 2016 to 2018 : Smile ECS, Toulouse
  - 2018 to now : Somfy Protect By Myfox, Toulouse
- Main work fields :
  - Complete linux firmwares for embedded systems
  - Connected alarm systems : video streaming, connected systems, embedded security

# Summary

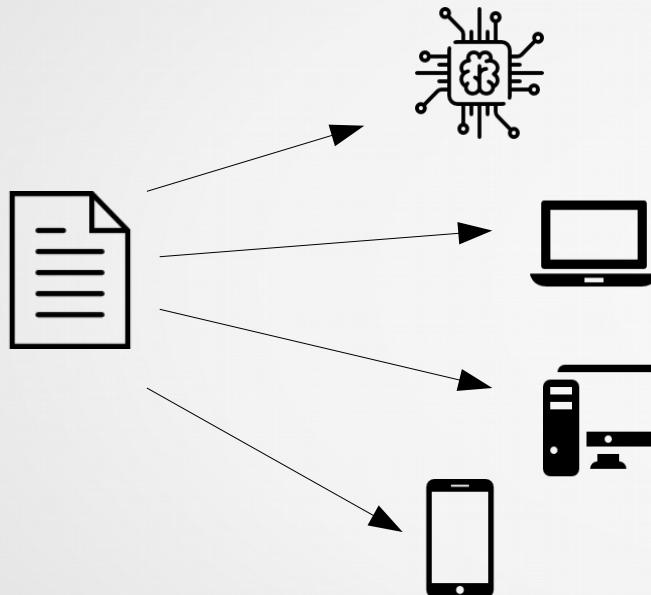
- What is V4L2
- Using V4L2 for video capture : initialization, stream loop
- Going further : encoding a raw image for image, for video

# What is V4L2

- A standardized kernel ↔ userspace API, started from kernel 2.1.X for v1, then dropped and replaced by v2 in v2.5.x
- Originally supports various captures support (webcams, TV tuners, radio)
- Now supports a large set of features : input/output for audio/video, tuners, modulators, extended user controls, M2M, etc.
- This presentation focuses on video capture



# Why should I use V4L2 ?



- One code base to rule them all !
- Can adapt to each platform at runtime (e.g. capabilities detection)
- But : you need proper V4L2 support (kernel and drivers)

# What do you need to begin

- Obviously, a video capture device
- A V4L2 capable linux-kernel
- Proper V4L2 driver for used device

```
[ 6877.094195] usb 1-6.2: new high-speed USB device number 11 using xhci_hcd
[ 6877.455368] usb 1-6.2: New USB device found, idVendor=046d, idProduct=0826, bcdDevice= 0.10
[ 6877.455372] usb 1-6.2: New USB device strings: Mfr=0, Product=2, SerialNumber=1
[ 6877.455375] usb 1-6.2: Product: HD Webcam C525
[ 6877.455377] usb 1-6.2: SerialNumber: B60A2F60
[ 6877.737575] usb 1-6.2: set resolution quirk: cval->res = 384
[ 6877.738303] uvcvideo: Found UVC 1.00 device HD Webcam C525 (046d:0826)
[ 6877.776277] uvcvideo 1-6.2:1.2: Entity type for entity Extension 5 was not initialized!
[ 6877.776280] uvcvideo 1-6.2:1.2: Entity type for entity Processing 2 was not initialized!
[ 6877.776281] uvcvideo 1-6.2:1.2: Entity type for entity Camera 1 was not initialized!
[ 6877.776283] uvcvideo 1-6.2:1.2: Entity type for entity Extension 6 was not initialized!
[ 6877.776285] uvcvideo 1-6.2:1.2: Entity type for entity Extension 7 was not initialized!
[ 6877.776286] uvcvideo 1-6.2:1.2: Entity type for entity Extension 8 was not initialized!
[ 6877.776353] input: HD Webcam C525 as /devices/pci0000:00/0000:00:14.0/usb1/1-6/1-6.2/1-6.2:1.2/input/input38
```

```
alexis@arch-alexis:~/|=> ls /dev/video*
/dev/video0  /dev/video1
```

# Your new friend : the official V4L2 doc

The screenshot shows a navigation sidebar on the left and a main content area on the right. The sidebar includes links for 'Introduction', 'Part I - Video for Linux API' (selected), 'Common API Elements', 'Image Formats', 'Input/Output', 'Interfaces', 'Libv4l Userspace Library', 'Changes', 'Function Reference', 'Common definitions for V4L2 and V4L2 subdev interfaces', 'Video For Linux Two Header File', 'Video Capture Example', and 'Video Grabber example'. The main content area has a header 'Part I - Video for Linux API' with a sub-header 'This part describes the Video for Linux API version 2 (V4L2 API) specification.' It also includes sections for 'Revision 4.5' and 'Table of Contents'.

Docs » Linux Media Infrastructure userspace API » Part I - Video for Linux API View page source

## Part I - Video for Linux API

This part describes the Video for Linux API version 2 (V4L2 API) specification.

### Revision 4.5

### Table of Contents

- 1. Common API Elements
  - 1.1. Opening and Closing Devices
    - 1.1.1. Device Naming
    - 1.1.2. Related Devices
    - 1.1.3. Multiple Opens
    - 1.1.4. Shared Data Streams
    - 1.1.5. Functions
  - 1.2. Querying Capabilities
  - 1.3. Application Priority
  - 1.4. Video Inputs and Outputs
    - 1.4.1. Example: Information about the current video input

- Can be used as an API documentation
- Can be used as a guide
- Userspace API, but also kernel API

The rest of the slides are heavily relying on this documentation

<https://linuxtv.org/downloads/v4l-dvb-apis/uapi/v4l/v4l2.html>

# Discovering your device

```
struct v4l2_capability cap ;  
  
dev_fd = open("/dev/video0", O_RDWR, 0) ;  
if(ioctl(dev_fd, VIDIOC_QUERYCAP, &cap) == -1)  
    fprintf(stderr, "Cannot fetch devices capabilities") ;  
  
if (!(cap.capabilities & V4L2_CAP_VIDEO_CAPTURE))  
    fprintf(stderr, "Selected device cannot capture video") ;
```

# Configuring your device input

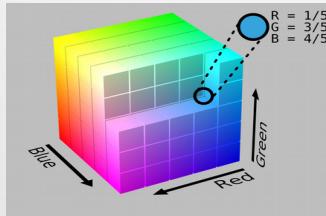
- setting the application priority : VIDIOC\_S\_PRIORITY
- querying / setting video input : VIDIOC\_ENUMINPUT, VIDIOC\_S\_INPUT
- querying / setting video standards : VIDIOC\_ENUMSTD
- cropping : VIDIOC\_S\_SELECTION
- setting streaming parameters VIDIOC\_S\_PARM
- and more...

## User controls on image

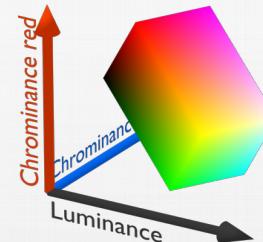
- A driver can expose many image controls (brightness, saturation, ...)
- We can enumerate and expose those controls : VIDIOC\_QUERYCTRL, VIDIOC\_QUERY\_EXT\_CTRL
- Some controls can be menus : VIDIOC\_QUERYMENU
- !\ VIDIOC\_QUERY\_EXT\_CTRL added later, check if supported by your kernel/driver !

# The pixel format : quick summary

- Many different image formats are available, depends on drivers :
  - different colorspaces : RGB, HSL, YCbCr...
  - single plane or multi plane format
  - indexed formats
  - etc
- Each format is represented with a FourCC code



RGB  
format



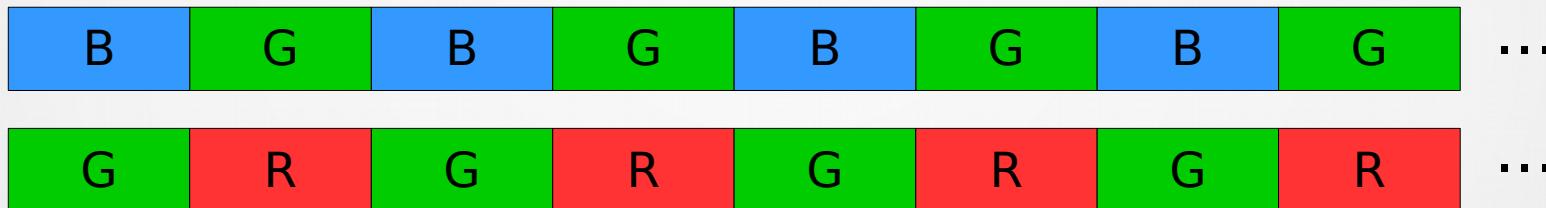
YCbCr  
format



HSL  
format

## Image format example : Bayer RGB

- FourCC code : 'B' 'A' '8' '1'
- Raw frame size : width \* height
- Why more G than B and R ? Because human eye is more sensible to green !



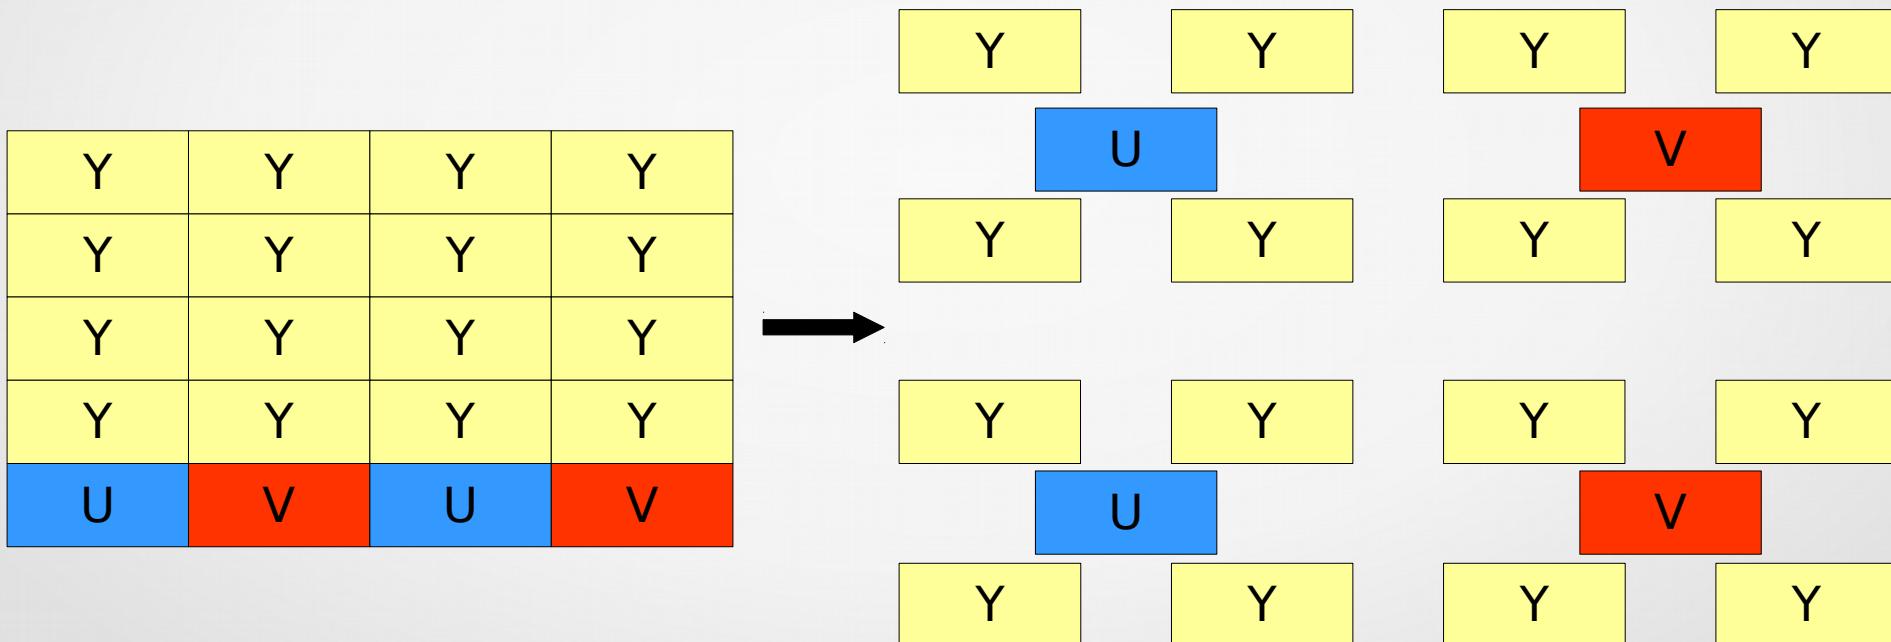
## Image format example : YUV422, or YUY2

- FourCC code : 'Y' 'U' 'Y' 'V'
- Raw frame size : width \* height \* 2
- Why more Y than U and V ? Because most of the image information is in Luma !



# Image format example : NV12, or YUV420SP

- FourCC code : 'N' 'V' '1' '2'
- Raw frame size : width \* height \* 3/2



# Negotiating image format

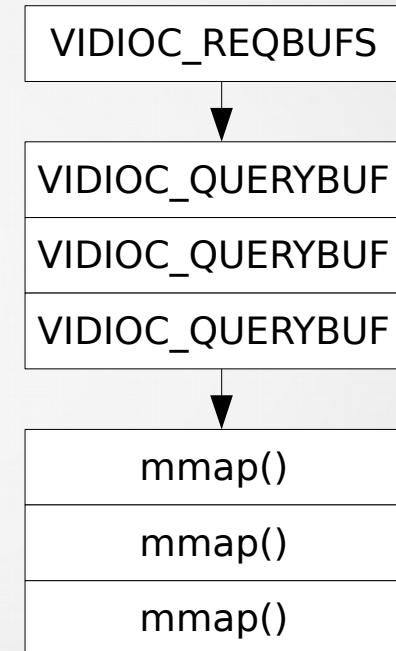
- Asking driver format capabilities : VIDIOC\_ENUM\_FMT
- Each format as a unique code, represented with 4 letters : FourCC code
- Selecting a format : VIDIOC\_S\_FMT
  - fill the v4l2\_format structure according to your needs
  - format is not expressed the same way between APIs (single vs multiplanar)
  - remember to check what the driver has really taken !

# Choosing a capture method

- Three way of retrieving raw frames :
  - basic I/O : read/write directly from/to device (remember to check for device capabilities : V4L2\_CAP\_READWRITE !)
  - streaming (needs V4L2\_CAP\_STREAMING !) :
    - memory mapping : request, mmap, enqueue and dequeue buffers
    - user pointers
    - DMA buffers (to export data between devices)

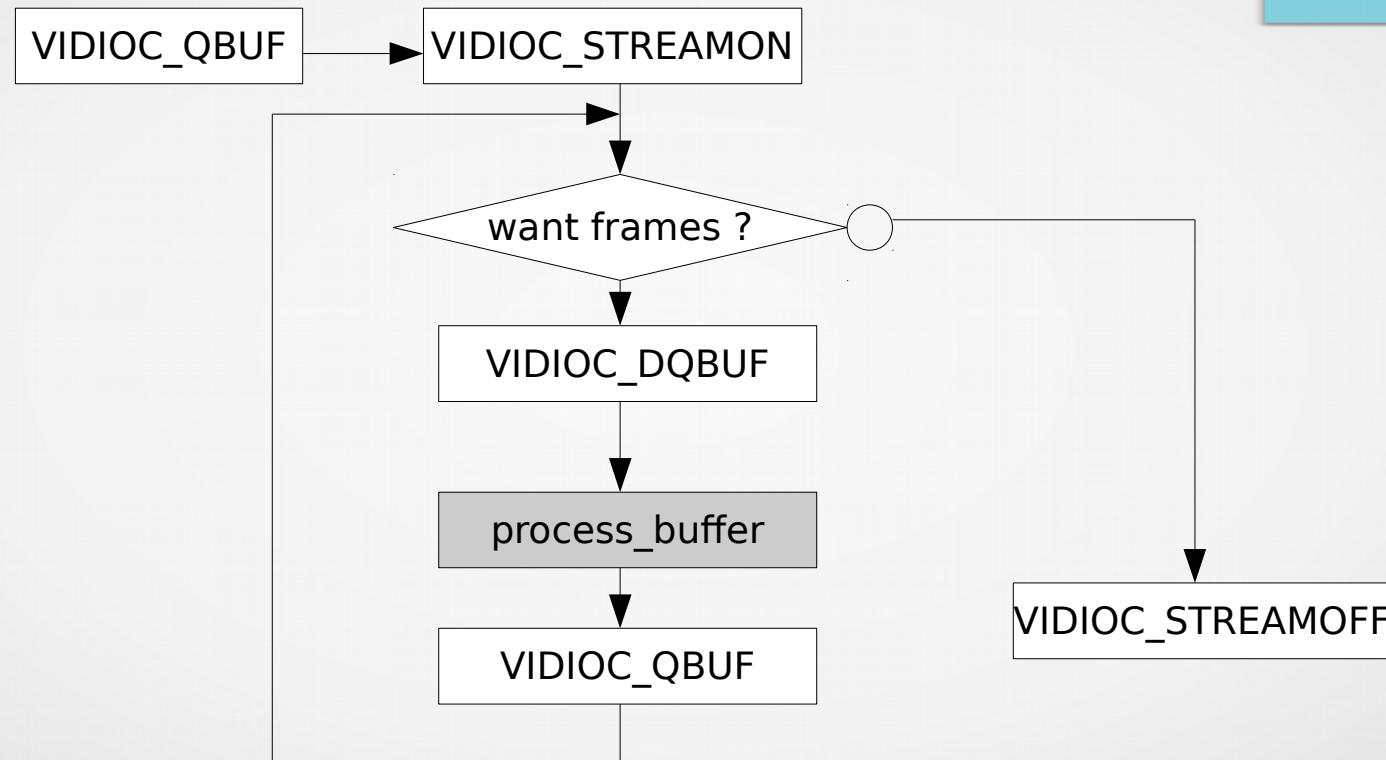
# Focusing on MMAP buffers

- Allocate buffers on driver side( set number of buffers, and type V4L2\_MEMORY\_MMAP)
- For each buffer, query length and location
- mmap the buffer



=> How many buffers do I need ???

# The capture loop



# Using raw images in real world ?

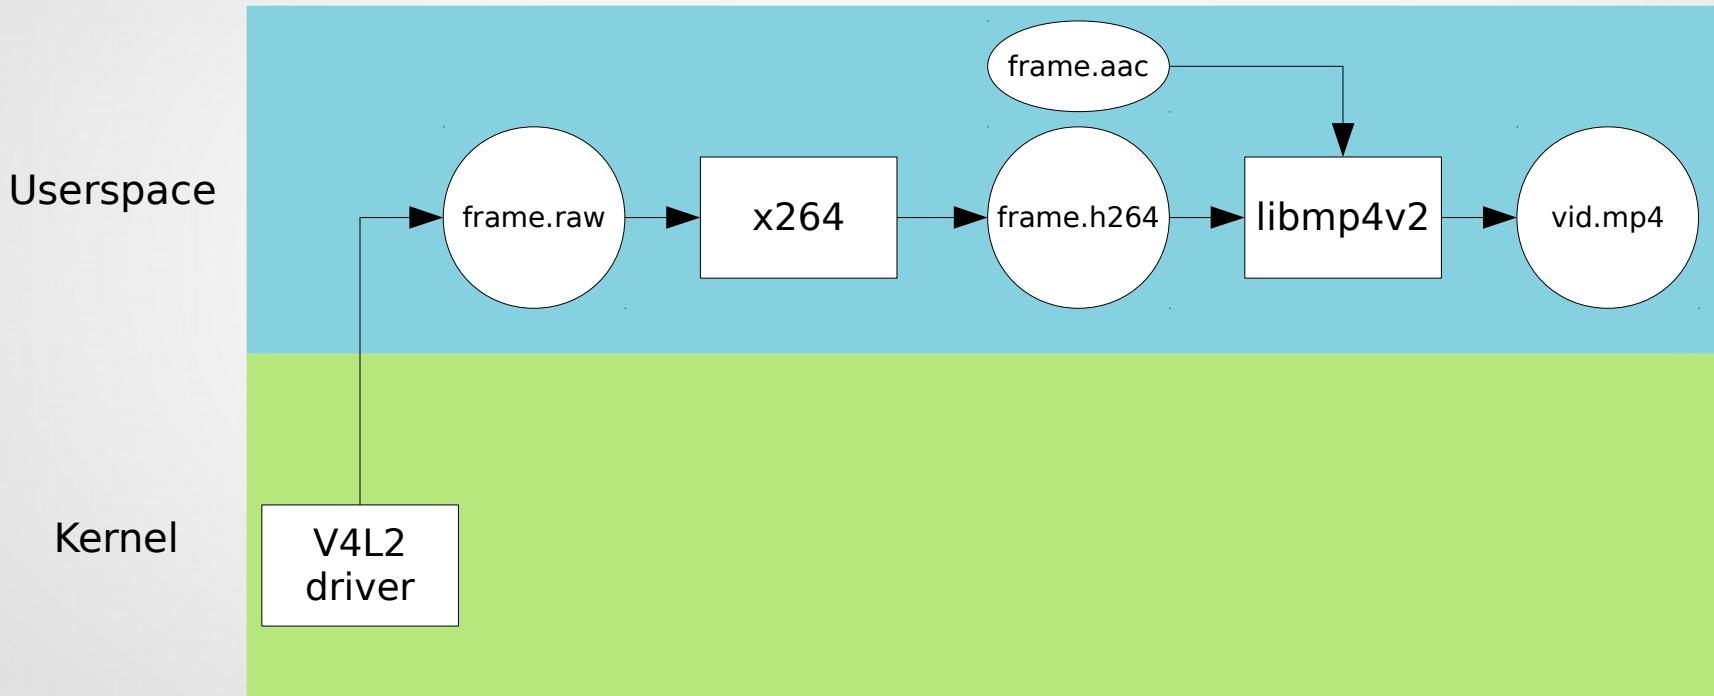
Displaying raw images

*<http://www.rawpixels.net/>*

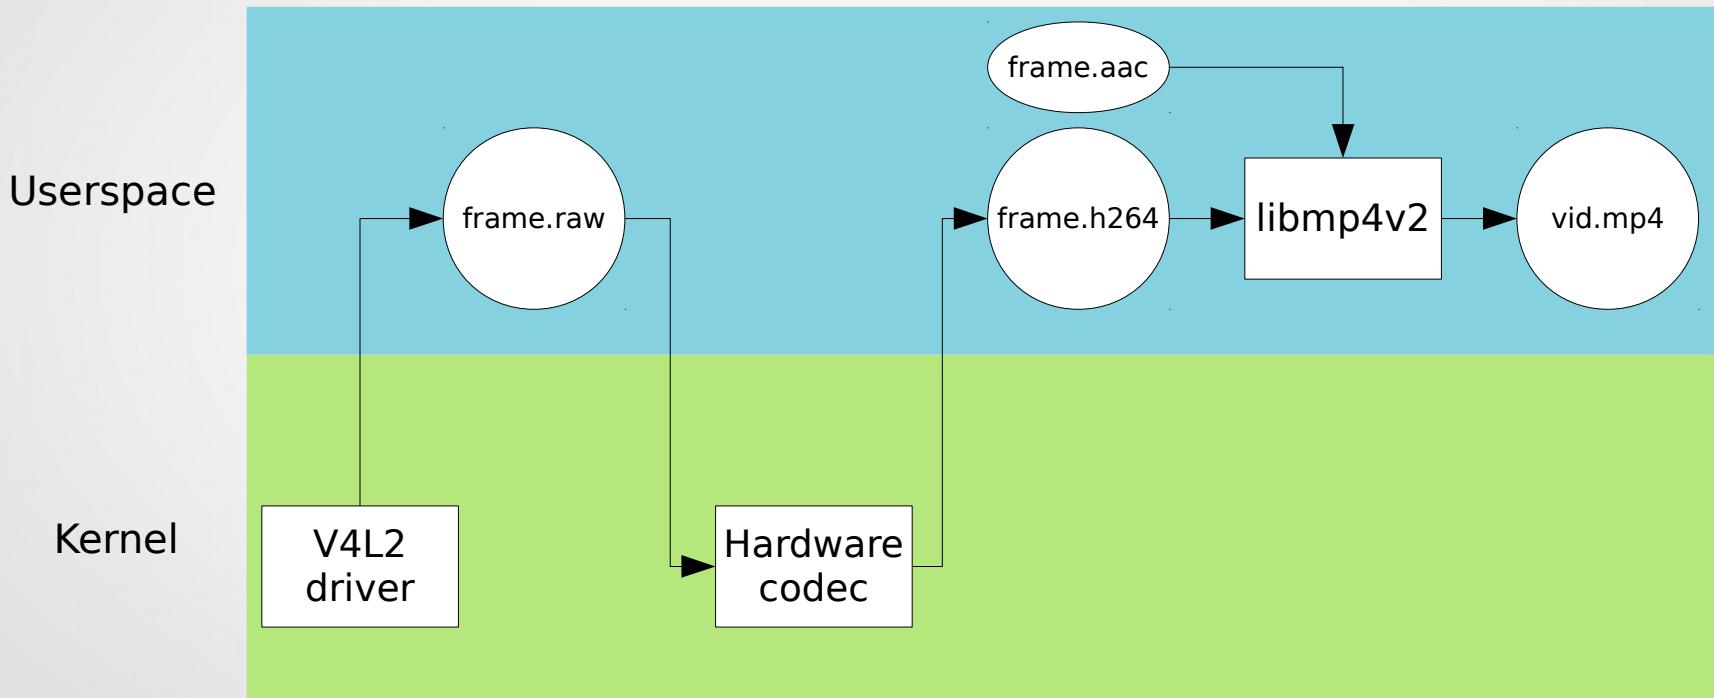
Encoding images :

we can pass raw RGB or YUV images to a JPEG encoder (software encoder, like jpeglib, or even a hardware JPEG encoder)

# Exploiting raw frames for video



# Exploiting raw frames for video



## Summary

- Let's face it : it was quite fastidious, just to get images... ( `cat src/*|wc -l` )  
968
- And we have just reinvented part of the wheel : see **v4l2-ctl**, **v4l2ucp**
- see also : ***libv4lconvert***, ***libv4l***, ***libv4l2***
- ***libcamera***, made to replace V4L2 ?

## Summary

- Having proper kernel support is sometimes a big assumption
  - Unimplemented ioctl ?
  - Unexpected behavior ?
- But we have a generic code which can run on multiple platforms !
- We have only discussed a very little part of the framework !

# Sources

- <https://fr.wikipedia.org/wiki/Video4Linux>
- <http://www.latelierducable.com/tv-televiseur/yuv-420-ycbcr-422-rgb-444-cest-quoi-le-chroma-subsampling/>
- <https://linuxtv.org/downloads/v4l-dvb-apis/uapi/v4l/v4l2.html>
- <http://linuxgizmos.com/libcamera-successor-to-v4l2-hopes-to-ease-embedded-linux-camera-headaches/>